



T3.1: Automating IoT device configuration. 1st June 2021

Renzo E. NAVAS (renzo.navas@inria.fr)

WP3: Virtualization and automation of IoT access networks.

Goals of Meeting

- Explain you what are we are doing in T3.1
- If possible, identify where the other partners (Acklio and Aguila?) can provide its expertise
 - (To Be Honest: we are relatively independent)

Recap T3.1: Goal

Description:

- Having a large number of IoT nodes to learn what parameters to use (time, power, spreading factor, etc.) for uploading data, and optimizing their global performance.
- Focus on mechanisms where IoT nodes make their own decisions (decentralized)
- But, end devices may have their strategy optimized globally by the orchestrator.

Goal:

- Propose and analyze the performance of machine learning algorithms that need few resources, like **multi-armed bandit methods**.
- **D3.1: Lightweight learning algorithms for massive IoT and analysis of their performance. (T15)**

Recap T3.1: Roadmap


Baseline: Do “better” than LoRaWAN’s ADR –and SoA–, in massive IoT scenarios.

1. Related Work (Positioning, Survey. E.g. [1]) [Non-priority for D3.1]
2. **Proposal:** Reinforcement-Learning-based, particularly **Bandit Algorithms** [2]
3. **Evaluation:** We will use a realistic setting/evaluation scenario, NS-3 based [URL].
4. **Contribution:** A differentiating factor of our Bandit-based Algorithm(s) will be this applicability/evaluation in realistic LoRaWAN scenarios.
 - In the literature, proposals use strong hypotheses or simplified models.
 - However, evaluation/comparison against non-bandit proposals will be a challenge (i.e., implementation)

Outline

- 01. Software Components (5min)
- 02. Bandits Recap (5min)
- 03. Bootstrapping LoRAWAN Bandits (~20)
- 04. Wrap Up

Software Components



*Core NS-3: A Discrete Event Simulator (in C++)

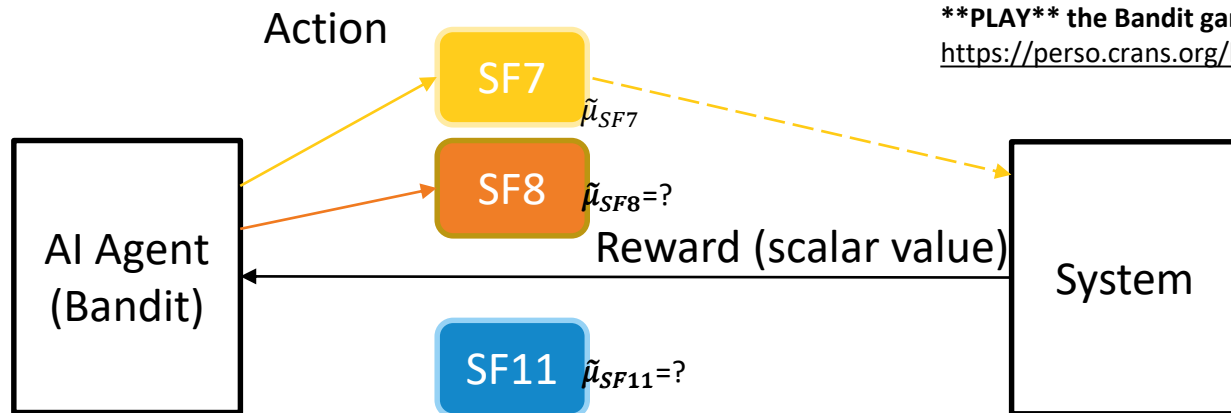
**LoRaWAN NS3 Module: <https://github.com/signetlabdei/lorawan>

AI-kernel: <https://github.com/Svalorzen/AI-Toolbox>

Bandit: Playing a game

****PLAY** the Bandit game for yourself! (Nice, Lilian B!):**

https://perso.crans.org/besson/phd/MAB_interactive_demo/



Objective:
Maximize the
long-run reward
(horizon N)



k Arms

Exploration: ~Refine estimation of Mean (μ) (and Variance!) of Arms
Exploitation: ~Use the arm with current max mean reward.

(most bandit techniques mix exploration with exploitation)

A bit of Detail (ADR Bandit Agent)

```
// @brief A class that implements a Bandit RL Agent targeted at Adaptive Data Rate LoRaWAN End Nodes
class AdrBanditAgent : public Object
{
public:
    /*...*/

    int ChooseArm ();
    //int chooseArmFromPolicyN(int policyNumber); If we have more than one policy! :O, all learn from shared experience

    void UpdateReward (int armNumber, double reward);

protected:
    Experience    m_experience;
    PolicyInterface * m_aiPolicy; // We can have more than one ML-Policy ☺, all can learn from the same Experience.
};
```


First Assumptions AI

- 1 Bandit Arm == 1 Spreading Factor
 - (Frequency will be uniformly random, and power constant)
 - 1 arm generalizes to anything (phy/mac params), but I want to keep total number of arms low (at least for now)
- Working on non-mobility Use Case (for now).
 - Mobility? Whatever we learn here can be adapted for mobility use cases (“first we learn to walk, then to fly”)
- Horizon infinite, BUT we want to converge relatively “fast” (i.e., Better than ADR)
 - In practical terms, we can assume we want to **converge for ~64 uplink packets**.
 - (Mobility UC will be equivalent to a finite horizon, with active triggering of a new learning phase/movement)

Adding Bandits to Sim Model - I

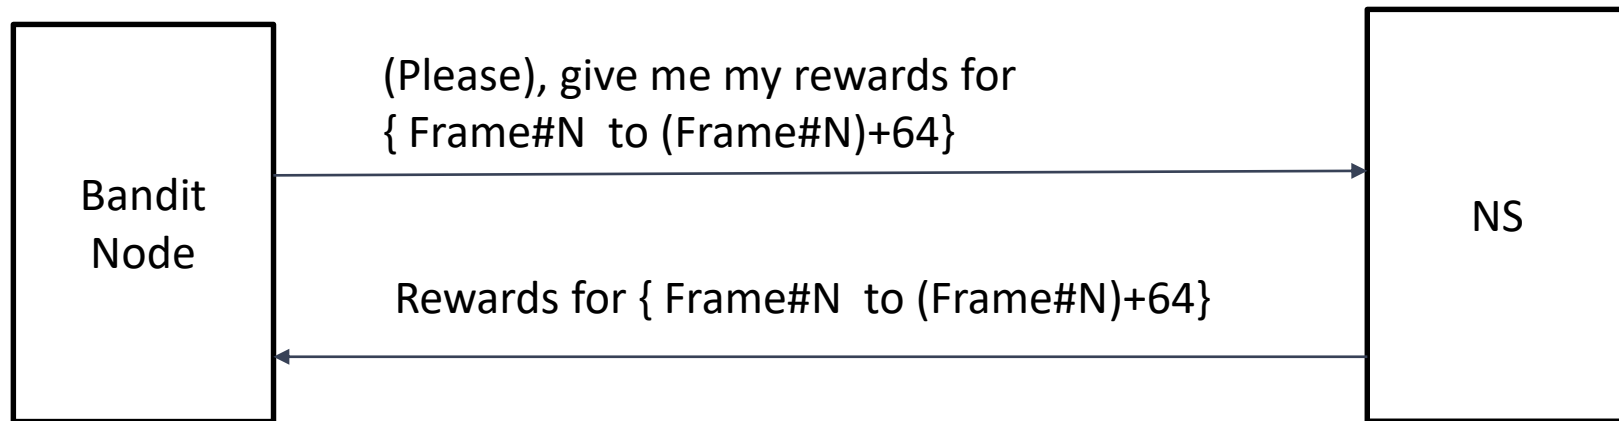
- A first working implem. of “Thompson Sampling” (TS) Bandit Policy
- New NS-3 Node: “Class A End Device LoRAWAN Mac Bandit”
- “Bandit” is an interface: use of any bandit is trivial.
 - Also, can be extended for more-than-bandits (e.g., Markov Decision Processes), with a system state var.
 - More-than-bandits is OK in sim, but unlikely on IoT nodes ☺ (not practical usability)
- Bandits C++ implementation from : <https://github.com/Svalorzen/AI-Toolbox>

Adding Bandits to Sim Model - II

- The (my) current challenge is the “Bootstrapping” of the policy (i.e., the first steps of the learning) → We want to converge “fast” ($\sim < 64$ msgs)
 - TS (and other policies, I think) needs **two** samples **per arm** to start (to have a **mean** and a **variance**).
 - Heuristically, we can do whatever we want to help bootstrapping (e.g., assume a priori values of mean and var.)
 - (*My problem...*) We know it is a waste of time to sample randomly/blindly all the arms in this bootstrapping phase. Because they are **not** independent in our LoRa use case... *we know relevant things* :)
 - *Another big challenge*: our feedback/reward is not immediate (DownLink is Expensive, and degenerates the medium performance); but, we can assume that in the “bootstrapping phase” we *actively request DLs to help converge fast* (the rate of request will be adaptive).

Delayed Feedback/Reward

- Remember, the Feedback message (either request or response) can be lost!
 - Ergo it can **not** be relative (e.g., “Give me last N packets rewards”): needs an **absolute reference** (i.e., Frame#)



Updates rewards for last 64 actions

Wrapping Up:

- A Functional TS Bandit in a LoRaWAN Simulation (Good!)
- But, early stages:
 - Refine the “bootstrapping” (including **REWARD** definition –fn(PDR,Energy)–)
 - Need to (define &) implement the “delayed feedback” MAC command
 - **Scalability** (N nodes, and proper way to gather Statistics)

Discussion

- Acklio, AGUILA involvement? (If it suits you)
- Sharing the Code (Private GitLab)
- Any Other?

Merci !

renzo.navas@inria.fr

Bibliography

- [1] Kufakunesu, Rachel, Gerhard P. Hancke, and Adnan M. Abu-Mahfouz. "A survey on Adaptive Data Rate optimization in LoRaWAN: Recent solutions and major challenges." Sensors 20.18 (2020): 5044.
- [2] Lattimore, Tor, and Csaba Szepesvári. Bandit algorithms. Cambridge University Press, 2020. URL: <https://tor-lattimore.com/downloads/book/book.pdf>

Slides from Previous Meeting

Bandit Algorithms – Introduction I



Round	1	2	3	4	5	6	7	8	9	10	11	
SF9	0		1		0	0				1	?	PDR = 2/5
SF7		1		0			0	0	0		?	PDR = 1/5

Arms ($k=2$) Exploration-phase Exploitation?

- Strategy: shall the IoT node keep using SF9, ignoring SF7?
- Or we attribute the poorer PDR performance of SF7 to “bad luck” and try it a few more times?
How many more times? (Exploration-Exploitation trade-off)
- ... What if we redefine the (negative) **reward** as the energy spent to deliver a packet?

Bandit Algorithms – Introduction II

- Goal: Maximize the cumulative reward, for a horizon of N “rounds” (could be infinite).
 - The definition of “reward” is fundamental (E.g., plain PDR vs Energy-aware PDR).
- In the basic bandit setting:
 - **Context-agnostic:** Arms’ rewards are independent (a reward for one arm does not give information about other arms’. E.g., a packet delivered with SF7 does not mean it also would have been delivered with SF9).
 - **Context-agnostic:** The learning agent **only** interacts with the system by “pulling” an arm and observing the empirical reward.
 - **Perfect Monitoring:** Empirical-reward observation (feedback) is immediate (... or almost).
 - **Stationary:** The underlying environment does not change (E.g., An arm’s reward’s “behavior” is always the same)

Bandit Algorithms – A realistic setting

- In our realistic setting:
 - **Contextual Bandits** → We can use contextual information (E.g., use DL for SNR stats, SFs are not-independent)
 - **Partial Monitoring** → We do not have perfect monitoring (E.g., DownLink is expensive)
 - **Non-stationary** → The environment, most likely, is not stationary.
- Defining an analytical model for this setting is not trivial, and theoretical bandit solutions does not exist (to the best of my knowledge).
- Our proposal(s) will use bandit solutions at their core, but probably will be mixed with **some heuristics**.
- Thanks to a realistic ns-3 evaluation/simulation environment, we will have strong **statistical guarantees** about their performance (and will compare with some SoA, including vanilla ADR). TBD: Obtain some theoretical guarantees.

LoRaWAN Evaluation: NS-3

- NS-3 LoRaWAN module
 - GIT : <https://github.com/signetlabdei/lorawan> and [Documentation](#).
 - Current version has an implementation of LoRaWAN's ADR.
 - Does not implement: Class-B Nodes, Frame Counters (but Lost packets are traced in-software).
 - Energy-measurement module could be implemented w/reasonable effort.
 - NS-3 module is highly customizable: E.g., stack of path loss models (Shadowing, Buildings..).

